

What could AI access earn your publication?

A directional estimate from your audience, your tokens,
and your share.

1 / 10

SWIPE →



Publisher AI revenue, explained.

AI assistants read publisher content to answer user questions.

Each query consumes a measurable number of tokens. Each token has a market price.

When publishers run the access layer, they keep a share of that revenue.

It is not a subscription. It is not an ad. It is a new line.

Two monetization patterns for publisher content.

SUBSCRIPTIONS & ADS

Priced per reader.

Capped by audience size.

Stable, slow growth.

PER-TOKEN AI ACCESS

Priced per token consumed.

Scales with AI usage.

New volume from agents.

Subscriptions price the reader. Per-token access prices the use.

Three inputs drive the estimate.

1

Audience

Monthly active users. Sessions per month. Questions per session.

2

Pricing

Tokens per query (2,000–4,000). Price per million tokens (\$20–\$150).

3

Share

The cut you keep from the per-query revenue.

How the math works.

$$\text{Revenue} = \text{questions} \times \text{tokens-per-query} \times \text{token-price} \times \text{your-share}$$

Audience and usage tell you how many queries reach your content each year.

Tokens and price tell you the gross value per query. Your share is what you keep after the deal.

WORKED EXAMPLE

$$10,000 \text{ users} \times 3 \text{ sessions} \times 1 \text{ question} \times 3,000 \text{ tokens} \times \$75 \text{ per million} \times 20\% \text{ share} = \$16,200 \text{ per year.}$$

Where the revenue actually comes from.

USER QUERY

An AI assistant asks a question on behalf of a user.



TOKEN CONSUMPTION

Tokens flow from publisher content into the answer.



AI PROVIDER

Charges the per-token rate to its end customer.

PUBLISHER

Keeps the agreed share of every query.

Per-token access is a new revenue line, not a re-slice of subscriptions.

What a realistic year looks like.

\$16,200 per year. \$1,350 a month. \$16 per active user.

Ten thousand monthly active users. Three sessions a month. One question per session. That is 360,000 questions a year, every one of them measured.

At 3,000 tokens per query, \$75 per million tokens, and a 20% publisher share, the model pencils out to a directional \$16,200.

FLOOR, NOT CEILING

These are modest assumptions. A larger catalog, a higher share, or a stronger deal moves the number up. Not down.

Audience scales the answer.

ACTIVE USERS	ANNUAL REVENUE
10,000	\$16,200
100,000	\$162,000
1,000,000	\$1.62M
5,000,000	\$8.1M

*Holding tokens at 3,000, price at \$75 per million, and share at 20%.
Revenue scales linearly with audience.*

What you can rely on. What you have to estimate.

RELIABLE

Active users

Sessions per month

Questions per session

DIRECTIONAL

Tokens per query (depends on retrieval depth)

Price per million tokens (depends on AI provider deal)

Your share (depends on the access contract)



Run the numbers on your catalog.

See what revenue your catalog could earn.

Visit cashmere.io