Cashmere

AI USAGE REPORT

Effective Downloads: a usage metric for the AI era

How publishers can measure content consumption when agents read in tokens, not items.

SWIPE →

DEFINITION

What is COUNTER?

COUNTER is the publisher industry standard for measuring content usage.

It tracks two item-level signals.

Investigations cover any interaction with an item or its metadata. Requests cover access to the content itself.

For decades, the ratio between them told publishers whether their content was being read, or just browsed.

COUNTER was built for a human usage pattern.

Search

↓

Investigate

↓

Request

↓

Read

One person. One item. One content decision.

The old system had two useful signals.

INVESTIGATION

Interaction with an item or its metadata.

Example: browsing the abstract.

Helped infer: interest / discovery.

REQUEST

Accessing the content itself.

Example: opening or downloading the item.

Helped infer: use / intent.

Together, they helped separate interest from use.

One AI query can become many retrieval events.

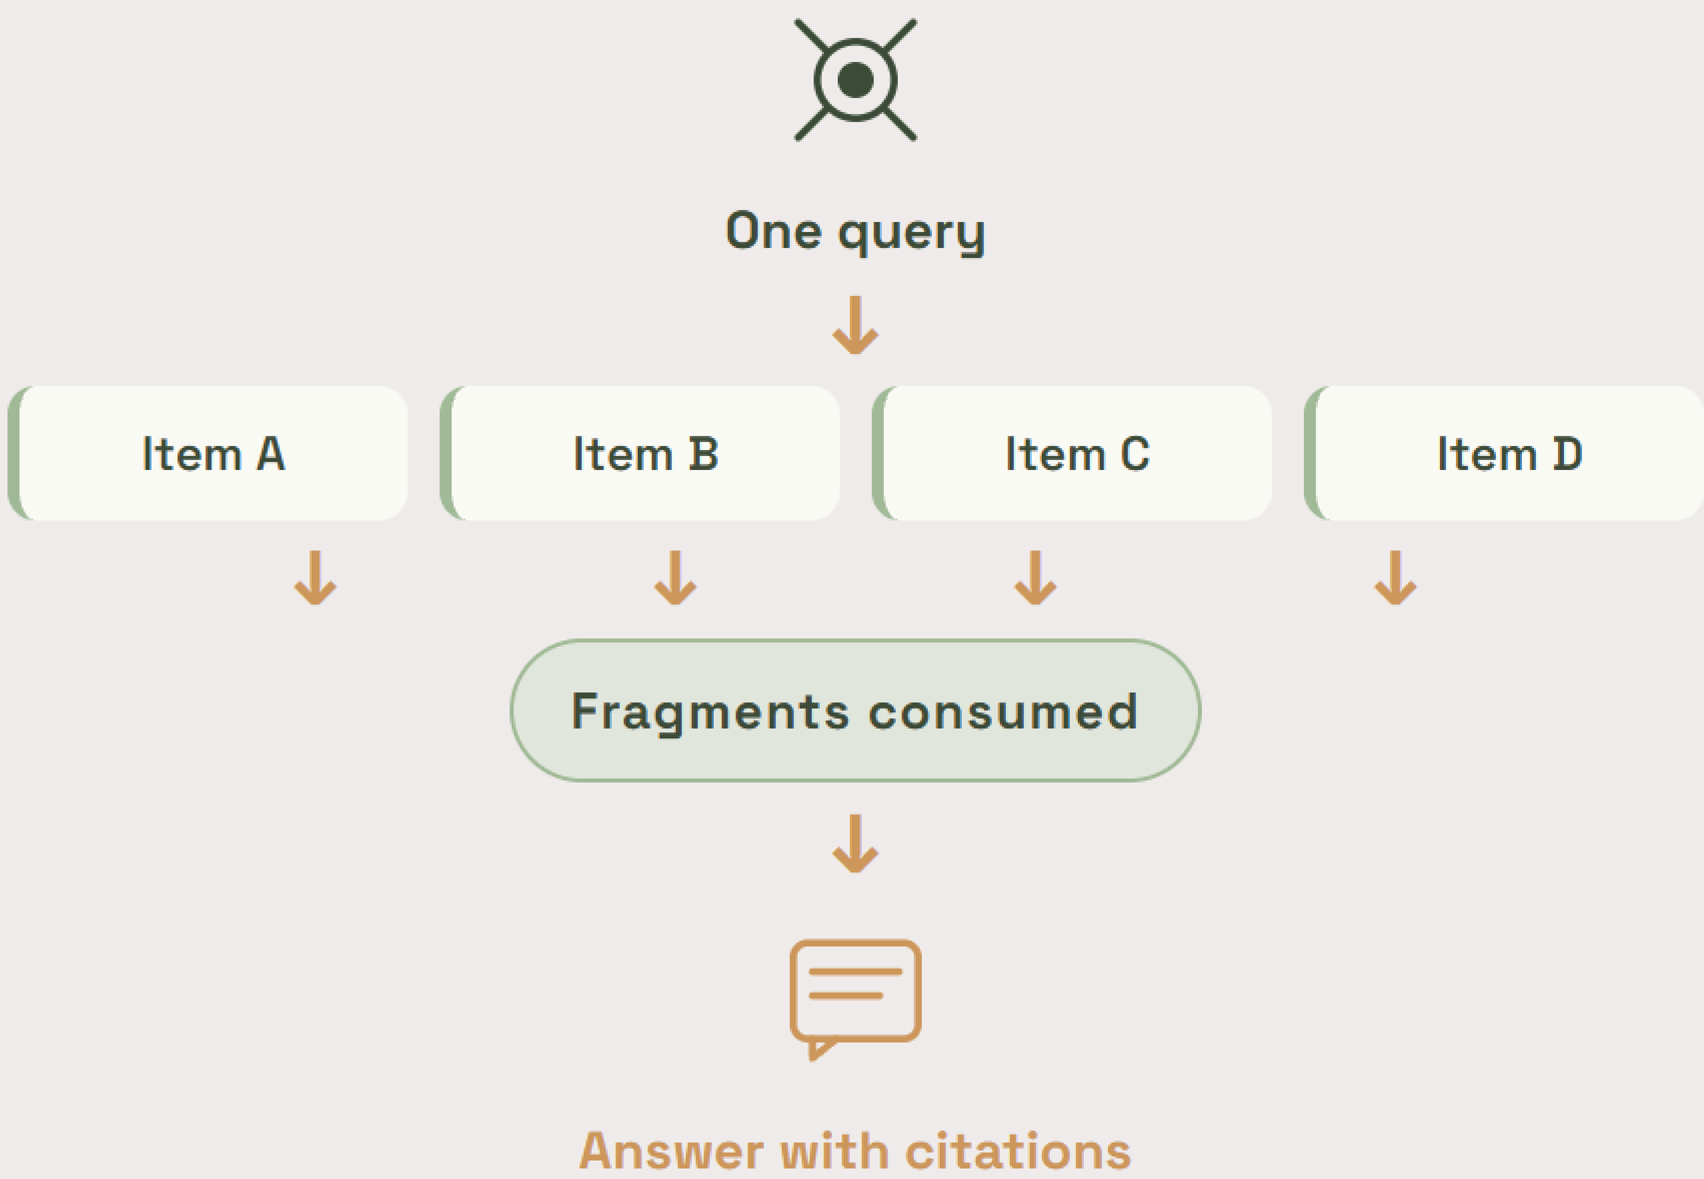

Agents may use pieces of many sources instead of reading one item whole.

The ratio starts measuring plumbing, not usage.

ACTUAL AGENT BEHAVIOUR

One query retrieves fragments from many items.

↓

PLATFORM PLUMBING

The access path decides what gets counted.

↙ ↘

REPORTED SIGNAL

Looks like browsing.

REPORTED SIGNAL

Looks like demand.

Same behaviour. Different reported reality.

A new unit: the Effective Download

> Effective Download = tokens consumed from an item ÷ total tokens in the item

AI agents don't read whole items. They pull a few thousand tokens — a couple of paragraphs, a figure caption, a conclusion — and move on.

Effective Downloads measure how much of an item an agent actually consumed. Partial reads, full reads, and repeated use are all captured by the same unit.

> EXAMPLE
>
> An article contains 10,000 tokens. An agent reads 5,000 of them — that's 0.5 Effective Downloads. If the same article gets referenced across many queries to 100,000 tokens of consumption, that's 10.

Tokens are stable where chunks are not.

CHUNKS LOW / INCONSISTENT	TOKENS HIGH / PORTABLE
Platform-specific.	Model-native.
Variable size.	Deterministic count.
Retrieval-dependent.	Fixed per item.
Not comparable across systems.	Comparable across AI systems.

> *That makes tokens a better unit for cross-platform usage measurement.*

Measurement has to happen upstream of the AI app.

AI APPLICATIONS

Chat, search, and research agents

Chat | Search | Research

↓

CASHMERE INFRASTRUCTURE

Token-level measurement layer

Retrievals | Tokens | Citations

↓

PUBLISHER USAGE INTELLIGENCE

Investigations, requests, Effective Downloads

Investigations | Requests | Effective Downloads

Cashmere

The metric is ready. The infrastructure is the missing layer.

Cashmere helps publishers make premium content AI-ready, governed, measurable, and monetizable.

Visit cashmere.io