

Anatomy of an AI Query

Three layers. One question, the content used to answer it, and the math underneath.

1 / 9

SWIPE →



Query • Chunk • Token

The query is whatever the user asks the AI. The chunks are the slices of content the AI pulls in to answer it. The tokens are the unit every answer gets billed in.

Once you have all three straight, the rest of how AI actually reads publisher content gets a lot easier to follow.

1. THE QUERY

Whatever the user asks.

A query is whatever the user asks the AI.

If you type a question, that's the query. If you upload a document and say "summarize this," the document plus your instruction is the query. Whatever input the user hands the model. That's the query.

Q Summarize this report in three|


Whatever the user asks – that's the query.

2. THE CHUNK

Three sources usually means three chunks.

A chunk is a small slice of a document, usually a few paragraphs. Publisher content gets indexed in chunks.

When an AI response says "based on three sources," it almost never means three articles. It means three chunks. Three paragraph-slices that matched the question.



CHUNKS

● Chunk 1 3 paragraphs

● Chunk 2 2 paragraphs

● Chunk 3 3 paragraphs

One article – three chunks. Each chunk can be retrieved on its own.

3. THE TOKEN

Text gets translated into tokens.

Before AI can process a chunk, the text gets translated. Roughly: one word, one token.

Each token has a numeric ID drawn from a vocabulary of about 200,000. A typical chunk is around 500 tokens.

TRY IT • LIVE TOKENIZER

AI doesn't read words – it reads tokens.

40 chars 9 tokens 4.44 chars/token


AI	·doesn't	·read	·words	·—	·it	·reads	·tokens	·
17527	8740	1729	6391	2733	480	31523	20290	13

Each colored chip is one token; the faint number below it is the token's numeric ID, which is what the model actually computes on. · marks a space.

How fast it adds up.

A single query pulls 5 to 100 chunks. At ~500 tokens per chunk, that's 2,500 to 50,000 tokens. For one question.

For reference, here is what those token counts feel like in print.



In Congress, July 4, 1776.

Declaration

Declaration of Independence
United States, 1776


TOKENS ~1,700



The New Yorker

A New Yorker feature
Long-form magazine
essay

TOKENS ~8,000



The Great GATSBY

The Great Gatsby
F. Scott Fitzgerald

TOKENS ~62,000



HARRY POTTER

and the Sorcerer's Stone

Harry Potter and the
Sorcerer's Stone
J. K. Rowling

TOKENS ~100,000


And then it scales.

Ask an AI agent: "What's the most popular restaurant in each of the 48 states?"

That looks like one question. To a modern AI agent, it is 48.

The agent fans out, splits the work, and runs a sub-query for each state. Each sub-agent pulls its own chunks. The math compounds.

Each sub-query is its own agent.



Each pulls its own chunks. The full task tallies up.



Tokens are the currency.

Queries enter, chunks participate, tokens get paid. Cashmere helps publishers measure and monetize each one.

Visit cashmere.io