

Pricing Content in the Age of AI

A proposed framework to better align publishers and AI applications.



Content Rate = Base Value × Use Case

Base Value captures the economic value embedded in the content itself. Use Case adjusts that value for the way AI consumes it.

One relationship that lets publishers and AI companies compare deals, benchmark outcomes, and negotiate from something other than instinct.

Base Value

The economic value in the content

Use Case

How the content is consumed

Every deal is judged in isolation.

Three publicly reported deals. No consistent way to tell if any of them was a good one.

1

~\$3,000 per book

Anthropic's proposed book settlement

2

\$250M+

News Corp's licensing agreement with OpenAI

3

~\$60M per year

Reddit's agreement with Google

DEFINITION

The Million Premium Token (MPT)

A label for tokens that carry deliberate human IP. Not a new kind of token. A unit of account, in the terms AI already prices in.

Music had the track. It carried the industry from the \$15 album to the \$0.99 download. Publishing never had one.

The token is the first unit that works across content types, business models, and AI applications.

EXAMPLE

Frontier labs already price output tokens at a 3 to 8x premium over input, because the value is in the intelligence behind them. Human content is the same: one string of tokens can be a quick social post, or years of research distilled.

The economic value in the content.

$$\text{Base Value} = Q \times \alpha \times U \times A$$

Four variables that translate a publisher's existing economics into a per-MPT rate. Not what the content is worth: a bestseller can be worth more than an article yet carry a lower dollar-per-MPT value. Distilled from analysis of more than 130 sources.

Q

The publisher's existing economic rate in its own channels

α

The share of that value attributable to the content itself

U

Uniqueness of the content relative to available alternatives

A

Relative authority, quality, and commercial position

Three ways AI consumes content.

The first step in pricing any AI use of content: identify which right is actually being licensed.

DIRECT RIGHTS

AI represents the original text back to a human

Reading, audio narration, translation

INFERENCE RIGHTS

AI consumes the text to generate something new

Grounding, deep research, summaries, agents

TRAINING RIGHTS

Original text improves the model itself

Out of scope for this brief

Two different rights. Two different economics.

DIRECT RIGHTS

An extension of rights
publishers already
license

Benchmark against
ebook, audio, and
translation economics

Use Case = Format
Multiplier

Text 1.0x. Add audio:
1.3 to 1.6x. Add
translation: 1.6 to 1.9x

INFERENCE RIGHTS

A fundamentally
different economic
model

One source text can
support millions of AI
interactions

Use Case =
Dependence × Market

Same text,
dramatically different
value by workflow

*Human consumption extends existing economics. AI consumption
needs new ones.*

The gap.

SCRAPER PRICE TODAY \$0.05 TO \$0.20 PER MPT, FLAT. PUBLISHER RECEIVES \$0	FRAMEWORK BASE VALUE ANCHORED TO MARKETS THAT ALREADY EXIST
Books	\$5 to \$540 per MPT
News & Media	\$130 to \$1,150 per MPT
Research & Reference	\$540 to \$8,310 per MPT

Underpriced by roughly 25x on books, up to 166,000x on professional reference. Not a disagreement over price. A different theory of value.

Will volume close the gap?

10,000x

more consumption needed

Calibrated against real usage from current AI applications and agentic workflows, today's scraper rates only make economic sense at roughly 10,000x current consumption. One day the industry gets there. It is not the immediate future.

SOURCE: PRICING CONTENT IN THE AGE OF AI, CASHMERE

Negotiate from math, not instinct.

The full framework, the assumptions behind it, and the worked examples are in the executive brief. The download link is in the post.

Visit cashmere.io